

CA3

Imports

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, IsolationForest
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

Reading data

```
In [17]: df = pd.read_csv("./assets/train.csv", index_col = 0)

df.info()

# Checking for missing values
print(f"\nMissing values inn training data: {df.isnull().sum().sum()}")
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2800 entries, -1.8257343 to -1.6260979
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Weight          2800 non-null  float64
1   Sweetness       2800 non-null  float64
2   Softness        2800 non-null  float64
3   HarvestTime     2800 non-null  float64
4   Ripeness        2800 non-null  float64
5   Acidity         2800 non-null  float64
6   Peel Thickness  2800 non-null  float64
7   Banana Density  2800 non-null  float64
8   Quality         2800 non-null  int64
dtypes: float64(8), int64(1)
memory usage: 218.8 KB
```

Missing values inn training data: 0

Data exploration and visualisation

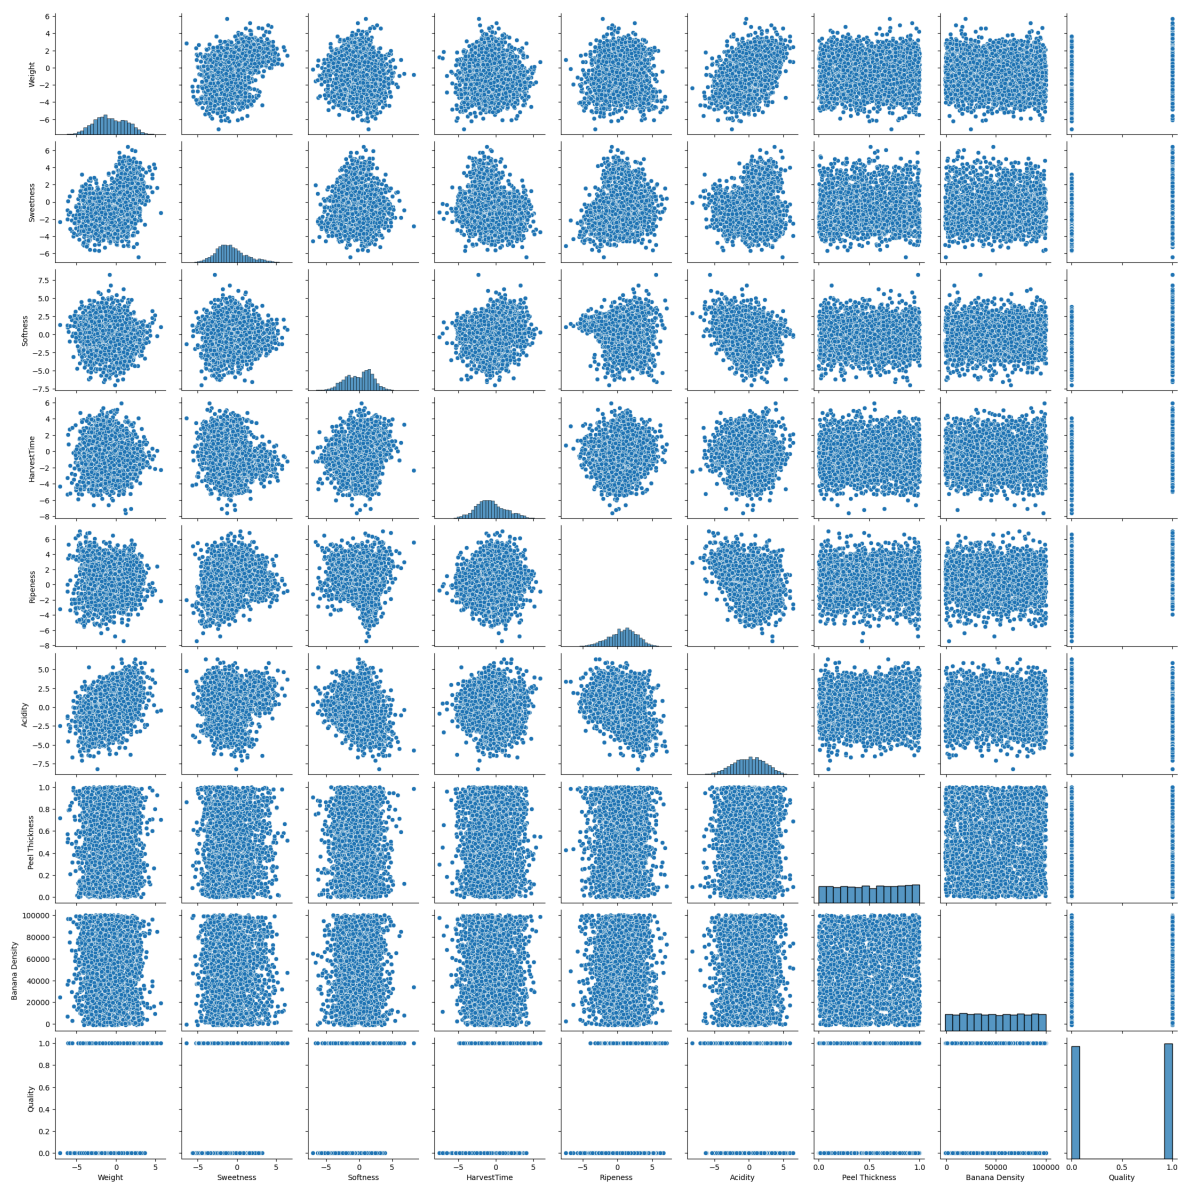
```
In [4]: print(df.describe()) # Gives a table of the dataset with statistical components

sns.pairplot(df) # Gives a scatter plot for every pair & a histogram for the the
```

	Weight	Sweetness	Softness	HarvestTime	Ripeness \
count	2800.000000	2800.000000	2800.000000	2800.000000	2800.000000
mean	-0.751050	-0.751005	-0.019557	-0.700683	0.771011
std	2.006590	1.955109	2.076865	2.029916	2.098275
min	-7.103426	-6.434022	-6.959320	-7.570008	-7.423155
25%	-2.238843	-2.104742	-1.593816	-2.112747	-0.572589
50%	-0.882387	-0.997902	0.220174	-0.856858	0.930927
75%	0.853566	0.334989	1.542899	0.628895	2.229410
max	5.679692	6.438196	8.241555	5.942060	7.077372

	Acidity	Peel Thickness	Banana Density	Quality
count	2800.000000	2800.000000	2800.000000	2800.000000
mean	-0.000989	0.506758	49397.491271	0.506429
std	2.286725	0.291936	29327.077623	0.500048
min	-8.226977	0.000086	-980.343999	0.000000
25%	-1.608385	0.257860	24025.427350	0.000000
50%	0.073963	0.506282	49303.534616	1.000000
75%	1.662417	0.761016	75066.598785	1.000000
max	6.395850	0.999430	99982.761410	1.000000

Out[4]: <seaborn.axisgrid.PairGrid at 0x2d288a22610>




The pairs of different features does not look linear separable, they have a complex relation.

Therefore, we think models like logistic regression will have low accuracy for these data.

```
In [5]: """ Checking correlation with every feature """
fig, axes = plt.subplots(1, 1, figsize=(14, 6)) # To make the values more readab
# annot = true to get value, coolwarm cmap to make sense with correlation
sns.heatmap(df.corr(), annot = True, cmap = "coolwarm")
```

Out[5]: <Axes: >



As we can see: Acidity and Weight, and Weight and Sweetness has the highest positive correlation.

While the Acidity and Ripeness, and Softness and Ripeness has the highest negative correlation.

To conclude:

The pairs of different features have complex relations to each other. The data is not linearly separable and

therefore, we think models like logistic regression will have low accuracy for these data.

Acidity and Weight, Weight and Sweetness has the highest positive correlation.

While the Acidity and Ripeness, and Softness and Ripeness has the highest negative correlation.

With this information we know which features we can focus more on.

We also found out that the banana density values are too high. We want to scale and standardize

the values of the dataset to fit the models better.

We will probably remove the columns banana density and peel thickness because of the low correlation

to the other features.

Data Preprocessing & Feature Engineering

We found out in the exploration and visualization part that there were not any missing values.

However, we will still drop outliers if we find any.

Feature Selection: We saw that Banana Density and Peel Thickness has low correlation with other features.

Therefore we will drop those features

```
In [19]: # Feature Selection, we are dropping the features when we're dropping columns i
useless_features = ["Banana Density", "Peel Thickness"]

# Including outliers, we found out that our model gets worse by removing outlier

df_clean = df.copy()
X = df_clean.drop(columns = ["Quality"] + useless_features)
y = df_clean["Quality"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, rand
```

Data preprocessing and visualisation

```
In [20]: # Standardizing
scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test) # Using transform to prevent data leak

print("Means:", X_train_scaled.mean(axis = 0))
print("\nStandard deviations:", X_train_scaled.std(axis = 0))

# Final Visualization
df_train_scaled = pd.DataFrame(X_train_scaled, columns = df_clean.drop(columns =


# Violin plots to visualize distribution
plt.figure(figsize = (12, 6))
sns.violinplot(data = df_train_scaled)
plt.xticks(rotation=60)
plt.show()

# Heatmap for visualizing correlation after data cleaning.
sns.heatmap(df_train_scaled.corr(), annot = True, cmap = "coolwarm")


# Pairplot for visualizing relations again.
sns.pairplot(df_train_scaled)
```


```
Means: [-1.94553368e-17  1.77635684e-17 -2.62224105e-17  1.69176842e-17
        -2.07241631e-17  0.00000000e+00]
```

```
Standard deviations: [1.  1.  1.  1.  1.  1.]
```



Out[20]: <seaborn.axisgrid.PairGrid at 0x2d2a5bc7610>





It looks pretty much the same. However the correlation has just changed a bit.
For example weight-acidity correlation changed with -0.01.

Modelling

Finding good parameters

```
In [9]: # Logistic Regression Classifier
param_grid_LR = {
    'C': [0.01, 0.1, 1, 10, 100],
}
grid_search_LR = GridSearchCV(LogisticRegression(), param_grid_LR, cv=5, n_jobs=
grid_search_LR.fit(X_train_scaled, y_train)

print(grid_search_LR.best_params_)

{'C': 1}
```

```
In [10]: # SVC
param_grid_SVC = {
    'C': [0.01, 0.1, 1, 10, 11],
    'kernel': ['linear', 'rbf', 'poly'],
    'gamma': ['scale', 'auto'],
}
```

```

}
grid_search_SVC = GridSearchCV(SVC(), param_grid_SVC, cv=5, n_jobs=-1)
grid_search_SVC.fit(X_train_scaled, y_train)

print(grid_search_SVC.best_params_)

```

```
{'C': 11, 'gamma': 'scale', 'kernel': 'rbf'}
```

```

In [11]: # DecisionTree
param_grid_DT = {
    'max_depth': [3, 10, 20, 30],
    'min_samples_split': [2, 5, 7, 8, 9, 10],
    'criterion': ['gini', 'entropy']
}
grid_search_DT = GridSearchCV(DecisionTreeClassifier(), param_grid_DT, cv=5, n_j
grid_search_DT.fit(X_train_scaled, y_train)

print(grid_search_DT.best_params_)

```

```
{'criterion': 'entropy', 'max_depth': 10, 'min_samples_split': 5}
```

```

In [12]: # Random Forest
param_grid_RF = {
    'n_estimators': [100, 125, 200],
    'max_depth': [15, 20, 21, 25],
    'criterion': ['gini', 'entropy']
}
grid_search_RF = GridSearchCV(RandomForestClassifier(), param_grid_RF, cv=5, n_j
grid_search_RF.fit(X_train_scaled, y_train)

print(grid_search_RF.best_params_)

```

```
{'criterion': 'gini', 'max_depth': 15, 'n_estimators': 200}
```

```

In [13]: # KNN
param_grid_KNN = {
    'n_neighbors': [3, 5, 7, 10, 15],
    'p': [1, 2, 3, 4, 5, 6, 7]
}
grid_search_KNN = GridSearchCV(KNeighborsClassifier(), param_grid_KNN, cv=5, n_j
grid_search_KNN.fit(X_train_scaled, y_train)

print(grid_search_KNN.best_params_)

```

```
{'n_neighbors': 5, 'p': 5}
```

Finding the best model

```

In [14]: dataset_sizes = np.arange(10, 2800, 50)

X_train_scaled = np.array(X_train_scaled)
y_train = np.array(y_train)

models = []

# Playing around with the variables to find something good
classifiers = {
    "LogisticRegression": LogisticRegression(C = 1),
    "SVC": SVC(C = 11, gamma = "scale", kernel = "rbf"),
    "DecisionTree": DecisionTreeClassifier(criterion = "entropy", max_depth = 10
    "RandomForest": RandomForestClassifier(criterion = "gini", max_depth = 20, n

```

```
    "KNN": KNeighborsClassifier(n_neighbors = 5, p = 5),
}

keys = classifiers.keys()

for clf_name in keys:
    clf = classifiers[clf_name]
    clf_index = list(keys).index(clf_name)

    for size_index, size in enumerate(dataset_sizes):
        X_train_subset, y_train_subset = X_train_scaled[:size], y_train[:size]

        clf.fit(X_train_subset, y_train_subset)

        y_pred = clf.predict(X_test_scaled)
        accuracy = accuracy_score(y_test, y_pred)

        models.append((clf, accuracy, size))

    print(f"{clf_name} | Size: {size} | Accuracy: {accuracy:.4f}")
```


LogisticRegression	Size: 10	Accuracy: 0.7857
LogisticRegression	Size: 60	Accuracy: 0.8600
LogisticRegression	Size: 110	Accuracy: 0.8671
LogisticRegression	Size: 160	Accuracy: 0.8643
LogisticRegression	Size: 210	Accuracy: 0.8614
LogisticRegression	Size: 260	Accuracy: 0.8600
LogisticRegression	Size: 310	Accuracy: 0.8643
LogisticRegression	Size: 360	Accuracy: 0.8657
LogisticRegression	Size: 410	Accuracy: 0.8600
LogisticRegression	Size: 460	Accuracy: 0.8600
LogisticRegression	Size: 510	Accuracy: 0.8600
LogisticRegression	Size: 560	Accuracy: 0.8586
LogisticRegression	Size: 610	Accuracy: 0.8614
LogisticRegression	Size: 660	Accuracy: 0.8657
LogisticRegression	Size: 710	Accuracy: 0.8629
LogisticRegression	Size: 760	Accuracy: 0.8657
LogisticRegression	Size: 810	Accuracy: 0.8657
LogisticRegression	Size: 860	Accuracy: 0.8700
LogisticRegression	Size: 910	Accuracy: 0.8700
LogisticRegression	Size: 960	Accuracy: 0.8671
LogisticRegression	Size: 1010	Accuracy: 0.8629
LogisticRegression	Size: 1060	Accuracy: 0.8614
LogisticRegression	Size: 1110	Accuracy: 0.8629
LogisticRegression	Size: 1160	Accuracy: 0.8657
LogisticRegression	Size: 1210	Accuracy: 0.8657
LogisticRegression	Size: 1260	Accuracy: 0.8671
LogisticRegression	Size: 1310	Accuracy: 0.8686
LogisticRegression	Size: 1360	Accuracy: 0.8671
LogisticRegression	Size: 1410	Accuracy: 0.8671
LogisticRegression	Size: 1460	Accuracy: 0.8686
LogisticRegression	Size: 1510	Accuracy: 0.8686
LogisticRegression	Size: 1560	Accuracy: 0.8671
LogisticRegression	Size: 1610	Accuracy: 0.8686
LogisticRegression	Size: 1660	Accuracy: 0.8700
LogisticRegression	Size: 1710	Accuracy: 0.8700
LogisticRegression	Size: 1760	Accuracy: 0.8700
LogisticRegression	Size: 1810	Accuracy: 0.8729
LogisticRegression	Size: 1860	Accuracy: 0.8729
LogisticRegression	Size: 1910	Accuracy: 0.8700
LogisticRegression	Size: 1960	Accuracy: 0.8729
LogisticRegression	Size: 2010	Accuracy: 0.8714
LogisticRegression	Size: 2060	Accuracy: 0.8729
LogisticRegression	Size: 2110	Accuracy: 0.8714
LogisticRegression	Size: 2160	Accuracy: 0.8714
LogisticRegression	Size: 2210	Accuracy: 0.8714
LogisticRegression	Size: 2260	Accuracy: 0.8714
LogisticRegression	Size: 2310	Accuracy: 0.8714
LogisticRegression	Size: 2360	Accuracy: 0.8714
LogisticRegression	Size: 2410	Accuracy: 0.8714
LogisticRegression	Size: 2460	Accuracy: 0.8714
LogisticRegression	Size: 2510	Accuracy: 0.8714
LogisticRegression	Size: 2560	Accuracy: 0.8714
LogisticRegression	Size: 2610	Accuracy: 0.8714
LogisticRegression	Size: 2660	Accuracy: 0.8714
LogisticRegression	Size: 2710	Accuracy: 0.8714
LogisticRegression	Size: 2760	Accuracy: 0.8714
SVC	Size: 10	Accuracy: 0.9000
SVC	Size: 60	Accuracy: 0.9171
SVC	Size: 110	Accuracy: 0.9200
SVC	Size: 160	Accuracy: 0.9357

SVC	Size: 210	Accuracy: 0.9529
SVC	Size: 260	Accuracy: 0.9543
SVC	Size: 310	Accuracy: 0.9500
SVC	Size: 360	Accuracy: 0.9514
SVC	Size: 410	Accuracy: 0.9571
SVC	Size: 460	Accuracy: 0.9586
SVC	Size: 510	Accuracy: 0.9657
SVC	Size: 560	Accuracy: 0.9629
SVC	Size: 610	Accuracy: 0.9671
SVC	Size: 660	Accuracy: 0.9657
SVC	Size: 710	Accuracy: 0.9643
SVC	Size: 760	Accuracy: 0.9686
SVC	Size: 810	Accuracy: 0.9700
SVC	Size: 860	Accuracy: 0.9700
SVC	Size: 910	Accuracy: 0.9671
SVC	Size: 960	Accuracy: 0.9657
SVC	Size: 1010	Accuracy: 0.9643
SVC	Size: 1060	Accuracy: 0.9671
SVC	Size: 1110	Accuracy: 0.9714
SVC	Size: 1160	Accuracy: 0.9686
SVC	Size: 1210	Accuracy: 0.9714
SVC	Size: 1260	Accuracy: 0.9714
SVC	Size: 1310	Accuracy: 0.9714
SVC	Size: 1360	Accuracy: 0.9729
SVC	Size: 1410	Accuracy: 0.9729
SVC	Size: 1460	Accuracy: 0.9729
SVC	Size: 1510	Accuracy: 0.9729
SVC	Size: 1560	Accuracy: 0.9714
SVC	Size: 1610	Accuracy: 0.9729
SVC	Size: 1660	Accuracy: 0.9729
SVC	Size: 1710	Accuracy: 0.9700
SVC	Size: 1760	Accuracy: 0.9700
SVC	Size: 1810	Accuracy: 0.9700
SVC	Size: 1860	Accuracy: 0.9714
SVC	Size: 1910	Accuracy: 0.9714
SVC	Size: 1960	Accuracy: 0.9700
SVC	Size: 2010	Accuracy: 0.9714
SVC	Size: 2060	Accuracy: 0.9714
SVC	Size: 2110	Accuracy: 0.9700
SVC	Size: 2160	Accuracy: 0.9700
SVC	Size: 2210	Accuracy: 0.9700
SVC	Size: 2260	Accuracy: 0.9700
SVC	Size: 2310	Accuracy: 0.9700
SVC	Size: 2360	Accuracy: 0.9700
SVC	Size: 2410	Accuracy: 0.9700
SVC	Size: 2460	Accuracy: 0.9700
SVC	Size: 2510	Accuracy: 0.9700
SVC	Size: 2560	Accuracy: 0.9700
SVC	Size: 2610	Accuracy: 0.9700
SVC	Size: 2660	Accuracy: 0.9700
SVC	Size: 2710	Accuracy: 0.9700
SVC	Size: 2760	Accuracy: 0.9700
DecisionTree	Size: 10	Accuracy: 0.5371
DecisionTree	Size: 60	Accuracy: 0.7886
DecisionTree	Size: 110	Accuracy: 0.7643
DecisionTree	Size: 160	Accuracy: 0.8114
DecisionTree	Size: 210	Accuracy: 0.8414
DecisionTree	Size: 260	Accuracy: 0.8529
DecisionTree	Size: 310	Accuracy: 0.8657
DecisionTree	Size: 360	Accuracy: 0.8900

DecisionTree	Size: 410	Accuracy: 0.8514
DecisionTree	Size: 460	Accuracy: 0.8714
DecisionTree	Size: 510	Accuracy: 0.8557
DecisionTree	Size: 560	Accuracy: 0.8971
DecisionTree	Size: 610	Accuracy: 0.8986
DecisionTree	Size: 660	Accuracy: 0.9029
DecisionTree	Size: 710	Accuracy: 0.9000
DecisionTree	Size: 760	Accuracy: 0.9129
DecisionTree	Size: 810	Accuracy: 0.9057
DecisionTree	Size: 860	Accuracy: 0.9114
DecisionTree	Size: 910	Accuracy: 0.9129
DecisionTree	Size: 960	Accuracy: 0.9271
DecisionTree	Size: 1010	Accuracy: 0.8971
DecisionTree	Size: 1060	Accuracy: 0.9114
DecisionTree	Size: 1110	Accuracy: 0.9057
DecisionTree	Size: 1160	Accuracy: 0.9129
DecisionTree	Size: 1210	Accuracy: 0.9100
DecisionTree	Size: 1260	Accuracy: 0.9143
DecisionTree	Size: 1310	Accuracy: 0.9000
DecisionTree	Size: 1360	Accuracy: 0.9143
DecisionTree	Size: 1410	Accuracy: 0.9086
DecisionTree	Size: 1460	Accuracy: 0.9129
DecisionTree	Size: 1510	Accuracy: 0.9129
DecisionTree	Size: 1560	Accuracy: 0.9100
DecisionTree	Size: 1610	Accuracy: 0.9200
DecisionTree	Size: 1660	Accuracy: 0.9086
DecisionTree	Size: 1710	Accuracy: 0.9129
DecisionTree	Size: 1760	Accuracy: 0.9171
DecisionTree	Size: 1810	Accuracy: 0.9129
DecisionTree	Size: 1860	Accuracy: 0.9143
DecisionTree	Size: 1910	Accuracy: 0.9286
DecisionTree	Size: 1960	Accuracy: 0.9314
DecisionTree	Size: 2010	Accuracy: 0.9229
DecisionTree	Size: 2060	Accuracy: 0.9300
DecisionTree	Size: 2110	Accuracy: 0.9229
DecisionTree	Size: 2160	Accuracy: 0.9229
DecisionTree	Size: 2210	Accuracy: 0.9200
DecisionTree	Size: 2260	Accuracy: 0.9200
DecisionTree	Size: 2310	Accuracy: 0.9214
DecisionTree	Size: 2360	Accuracy: 0.9243
DecisionTree	Size: 2410	Accuracy: 0.9229
DecisionTree	Size: 2460	Accuracy: 0.9186
DecisionTree	Size: 2510	Accuracy: 0.9271
DecisionTree	Size: 2560	Accuracy: 0.9243
DecisionTree	Size: 2610	Accuracy: 0.9243
DecisionTree	Size: 2660	Accuracy: 0.9214
DecisionTree	Size: 2710	Accuracy: 0.9257
DecisionTree	Size: 2760	Accuracy: 0.9200
RandomForest	Size: 10	Accuracy: 0.7586
RandomForest	Size: 60	Accuracy: 0.8743
RandomForest	Size: 110	Accuracy: 0.8671
RandomForest	Size: 160	Accuracy: 0.8986
RandomForest	Size: 210	Accuracy: 0.9100
RandomForest	Size: 260	Accuracy: 0.9200
RandomForest	Size: 310	Accuracy: 0.9314
RandomForest	Size: 360	Accuracy: 0.9443
RandomForest	Size: 410	Accuracy: 0.9386
RandomForest	Size: 460	Accuracy: 0.9471
RandomForest	Size: 510	Accuracy: 0.9457
RandomForest	Size: 560	Accuracy: 0.9471

RandomForest	Size: 610	Accuracy: 0.9586
RandomForest	Size: 660	Accuracy: 0.9514
RandomForest	Size: 710	Accuracy: 0.9486
RandomForest	Size: 760	Accuracy: 0.9571
RandomForest	Size: 810	Accuracy: 0.9529
RandomForest	Size: 860	Accuracy: 0.9543
RandomForest	Size: 910	Accuracy: 0.9614
RandomForest	Size: 960	Accuracy: 0.9557
RandomForest	Size: 1010	Accuracy: 0.9557
RandomForest	Size: 1060	Accuracy: 0.9557
RandomForest	Size: 1110	Accuracy: 0.9543
RandomForest	Size: 1160	Accuracy: 0.9586
RandomForest	Size: 1210	Accuracy: 0.9557
RandomForest	Size: 1260	Accuracy: 0.9557
RandomForest	Size: 1310	Accuracy: 0.9543
RandomForest	Size: 1360	Accuracy: 0.9557
RandomForest	Size: 1410	Accuracy: 0.9543
RandomForest	Size: 1460	Accuracy: 0.9529
RandomForest	Size: 1510	Accuracy: 0.9571
RandomForest	Size: 1560	Accuracy: 0.9529
RandomForest	Size: 1610	Accuracy: 0.9571
RandomForest	Size: 1660	Accuracy: 0.9543
RandomForest	Size: 1710	Accuracy: 0.9529
RandomForest	Size: 1760	Accuracy: 0.9571
RandomForest	Size: 1810	Accuracy: 0.9571
RandomForest	Size: 1860	Accuracy: 0.9600
RandomForest	Size: 1910	Accuracy: 0.9557
RandomForest	Size: 1960	Accuracy: 0.9543
RandomForest	Size: 2010	Accuracy: 0.9586
RandomForest	Size: 2060	Accuracy: 0.9543
RandomForest	Size: 2110	Accuracy: 0.9557
RandomForest	Size: 2160	Accuracy: 0.9586
RandomForest	Size: 2210	Accuracy: 0.9571
RandomForest	Size: 2260	Accuracy: 0.9571
RandomForest	Size: 2310	Accuracy: 0.9557
RandomForest	Size: 2360	Accuracy: 0.9586
RandomForest	Size: 2410	Accuracy: 0.9557
RandomForest	Size: 2460	Accuracy: 0.9557
RandomForest	Size: 2510	Accuracy: 0.9586
RandomForest	Size: 2560	Accuracy: 0.9600
RandomForest	Size: 2610	Accuracy: 0.9600
RandomForest	Size: 2660	Accuracy: 0.9571
RandomForest	Size: 2710	Accuracy: 0.9586
RandomForest	Size: 2760	Accuracy: 0.9600
KNN	Size: 10	Accuracy: 0.6429
KNN	Size: 60	Accuracy: 0.9243
KNN	Size: 110	Accuracy: 0.9386
KNN	Size: 160	Accuracy: 0.9486
KNN	Size: 210	Accuracy: 0.9514
KNN	Size: 260	Accuracy: 0.9557
KNN	Size: 310	Accuracy: 0.9557
KNN	Size: 360	Accuracy: 0.9514
KNN	Size: 410	Accuracy: 0.9571
KNN	Size: 460	Accuracy: 0.9586
KNN	Size: 510	Accuracy: 0.9571
KNN	Size: 560	Accuracy: 0.9586
KNN	Size: 610	Accuracy: 0.9571
KNN	Size: 660	Accuracy: 0.9600
KNN	Size: 710	Accuracy: 0.9600
KNN	Size: 760	Accuracy: 0.9614

KNN	Size: 810	Accuracy: 0.9586
KNN	Size: 860	Accuracy: 0.9586
KNN	Size: 910	Accuracy: 0.9571
KNN	Size: 960	Accuracy: 0.9571
KNN	Size: 1010	Accuracy: 0.9571
KNN	Size: 1060	Accuracy: 0.9586
KNN	Size: 1110	Accuracy: 0.9614
KNN	Size: 1160	Accuracy: 0.9629
KNN	Size: 1210	Accuracy: 0.9643
KNN	Size: 1260	Accuracy: 0.9629
KNN	Size: 1310	Accuracy: 0.9629
KNN	Size: 1360	Accuracy: 0.9586
KNN	Size: 1410	Accuracy: 0.9600
KNN	Size: 1460	Accuracy: 0.9614
KNN	Size: 1510	Accuracy: 0.9614
KNN	Size: 1560	Accuracy: 0.9614
KNN	Size: 1610	Accuracy: 0.9643
KNN	Size: 1660	Accuracy: 0.9629
KNN	Size: 1710	Accuracy: 0.9614
KNN	Size: 1760	Accuracy: 0.9586
KNN	Size: 1810	Accuracy: 0.9586
KNN	Size: 1860	Accuracy: 0.9600
KNN	Size: 1910	Accuracy: 0.9586
KNN	Size: 1960	Accuracy: 0.9586
KNN	Size: 2010	Accuracy: 0.9586
KNN	Size: 2060	Accuracy: 0.9586
KNN	Size: 2110	Accuracy: 0.9586
KNN	Size: 2160	Accuracy: 0.9586
KNN	Size: 2210	Accuracy: 0.9586
KNN	Size: 2260	Accuracy: 0.9586
KNN	Size: 2310	Accuracy: 0.9586
KNN	Size: 2360	Accuracy: 0.9586
KNN	Size: 2410	Accuracy: 0.9586
KNN	Size: 2460	Accuracy: 0.9586
KNN	Size: 2510	Accuracy: 0.9586
KNN	Size: 2560	Accuracy: 0.9586
KNN	Size: 2610	Accuracy: 0.9586
KNN	Size: 2660	Accuracy: 0.9586
KNN	Size: 2710	Accuracy: 0.9586
KNN	Size: 2760	Accuracy: 0.9586

Final evaluation

```
In [15]: # Get the model with best accuracy by choosing the maximum index 1 in the nested
best_model = max(models, key = lambda x: x[1])
print(f"{best_model[0]} | Size: {best_model[2]} | Accuracy: {best_model[1]}")
```

SVC(C=11) | Size: 1360 | Accuracy: 0.9728571428571429

This is our best Model:

SVC with C = 11

Size on dataset 1360

With an Accuracy 97.29%,

The accuracy will get higher on the kaggle submission.

Kaggle submission

```
In [16]: # Our best model.
model = best_model[0]

# Getting our test data and dropping useless features.
df_test = pd.read_csv("./assets/test.csv", index_col = 0)
df_test = df_test.drop(columns = useless_features)

# Scaling
X_test2_scaled = scaler.transform(df_test)

# Taken from the kaggle submission side.
y_test2 = model.predict(X_test2_scaled)
y_test2 = pd.DataFrame(y_test2, columns=["Quality"])
y_test2.index.name = "ID"
y_test2[['Quality']].to_csv("submission.csv")
```